

# **Mastering ELK + EFK**

# Elasticsearch

- A search engine based on the Lucene library.
- It provides a distributed, multitenant-capable full-text search engine with an HTTP web interface and schema-free JSON documents.
- Developed in Java



# Elasticsearch is Document Oriented

- **Insert** Documents
- **Delete** Documents
- **Retrieve** Documents
- **Analyze** Documents
- **Search** Documents

## Elasticsearch is Document Oriented


```
{  
  "total": 6,  
  "Type": "employees",  
  "documents": [  
    {  
      "name": "David",  
      "id": 1101,  
      "salary": 41000,  
      "hiredate": "October 3, 1998",  
      "department": "admin"  
    },  
    {  
      "name": "Michelle",  
      "id": 1103,  
      "salary": 45000,  
      "hiredate": "April 7, 2008",  
      "department": "Research"  
    },  
    {  
      "name": "Cassandra",  
      "id": 1102,  
      "salary": 68000,  
      "hiredate": "January 12, 2001",  
      "department": "Sales"  
    },  
    {  
      "name": "Brian",  
      "id": 1104,  
      "salary": 37000,  
      "hiredate": "August 19, 2012",  
      "department": "Admin"  
    },  
    {  
      "name": "Jason",  
      "id": 1105,  
      "salary": 92000,  
      "hiredate": "March 15, 2013",  
      "department": "Sales"  
    },  
    {  
      "name": "Robert",  
      "id": 1106,  
      "salary": 43000,  
      "hiredate": "January 11, 2014",  
      "department": "Sales"  
    }  
  ]  
}
```



# General Features

- Scalable up to petabytes of structured and unstructured data.
- Can be used as a replacement of document stores like MongoDB.
- Uses denormalization to improve the search performance.
- Is one of the popular enterprise search engines
- Is currently being used by many big organizations like
  - Wikipedia,
  - The Guardian,
  - StackOverflow,
  - GitHub etc.
- Is an open source and available under the Apache license version 2.0.

# Key Concepts




# Key Concepts

- Node
  - Single running instance of Elasticsearch.
- Cluster
  - Collection of one or more nodes.
  - Cluster provides collective indexing and search capabilities across all the nodes for entire data.
- Index
  - Collection of different type of documents and their properties.
  - Uses the concept of shards to improve the performance.
- Document
  - Collection of fields in a specific manner defined in JSON format.


# Key Concepts

- Shard
  - Indexes are horizontally subdivided into shards.
  - This means each shard contains all the properties of document but contains less number of JSON objects.
  - Primary shard is the original horizontal part of an index and then these primary shards are replicated into replica shards



# Key Concepts

- Replicas
  - Replicas of indexes and shards.
  - Helps in increasing the availability of data
  - Improves the performance of searching
    - Parallel search operation in these replicas.



# Index > Type > Document > Field

- **Vehicles (index)**


- **Cars (type)**

- Car (document 1)
    - Car (document 2)
    - Car (document n)




- **Motorcycles (Type)**

- Motor (document 1)
    - Motor (document 2)
    - Motor (document n)



# Index Request



# Inserting = Indexing


# Advantages

- Developed on Java
  - Makes it compatible on almost every platform.
- Real time
  - After one second the added document is searchable in this engine.
- Distributed
  - Makes it easy to scale and integrate in any big organization.
- Uses JSON objects
  - Makes it possible to invoke the Elasticsearch server with a large number of different programming languages.

# Disadvantages

- Does not have multi-language support in terms of handling request and response data (only possible in JSON)
- Security
  - Does not provide any built-in authentication or access control functionality.
- Transactions
  - There is no much more support for transactions or processing on data manipulation.

# Comparison between Elasticsearch and RDBMS



# Syntax for indexing a document

```
PUT /{index}/{type}/{id}
{
    "field1": "value1",
    "field2": "value2",
    ...
}
```

# Data Type for Document Fields

**String Fields:** text, keyword

**Numeric Fields:** long, integer, short, byte, double, float

**Date Fields:** text, keyword

**True/False Fields:** boolean

**Binary Fields:** binary

# Inverted Index

Maps words to the actual document locations of where they occur



[lucene.apache.org](http://lucene.apache.org)

# Inverted Index

- 1: Winter is coming.
- 2: Ours is the fury.
- 3: The choice is yours.

<u>term</u>	<u>freq</u>	<u>documents</u>
choice	1	3
coming	1	1
fury	1	2
is	3	1, 2, 3
ours	1	2
the	2	2, 3
winter	1	1
yours	1	3

Dictionary      Postings

**Document 1**

The bright blue butterfly hangs on the breeze.

**Document 2**

It's best to forget the great sky and to retire from every wind.

**Document 3**

Under blue sky, in bright sunlight, one need not search around.

**Stopword list**

a  
and  
around  
every  
for  
from  
in  
is  
it  
not  
on  
one  
the  
to  
under

**Inverted index**


ID	Term	Document
1	best	2
2	blue	1, 3
3	bright	1, 3
4	butterfly	1
5	breeze	1
6	forget	2
7	great	2
8	hangs	1
9	need	3
10	retire	2
11	search	3
12	sky	2, 3
13	wind	2

# Logstash

- Logstash is a tool for managing events and logs.
- Encompasses a larger system of log collection, processing, storage and searching activities.
- Part of the Elastic Stack along with Beats, Elasticsearch and Kibana
- Processing pipeline that
  - Ingests data from a multitude of sources simultaneously,
  - Transforms it, and then
  - Sends it to Elasticsearch
- Logstash has over 200 plugins, and you can write your own very easily as well.



# Logstash



# Elastic Search – Populate


- Can use the following command to create an index:
  - PUT school
- Response
  - {"acknowledged": true}

# Elastic Search – Populate


- Add data
  - POST school/\_doc/10
  - {
  - "name":"Saint Paul School", "description":"ICSE Afiliation",
  - "street":"Dawarka", "city":"Delhi", "state":"Delhi", "zip":"110075",
  - "location": [28.5733056, 77.0122136], "fees":5000,
  - "tags": ["Good Faculty", "Great Sports"], "rating": "4.5"
  - }

# Adding Sample Data in Kibana


- Kibana is a GUI driven tool for accessing the data and creating the visualization.
- In the Kibana home page, choose the following option to add sample ecommerce data:




# Search Request




# Elastic Stack ecosystem




# A JSON representation of a book entity




# Elasticsearch URL invocation endpoint




# Shards of data distributed across nodes



# Accessing the DevTools navigation page



# Elasticsearch's request and response flow



# Indexing two more documents using document API

```
PUT books/_doc/2
{
  "title": "Core Java Volume I - Fundamentals",
  "author": "Cay S. Horstmann",
  "release_date": "2018-08-27",
  "amazon_rating": 4.8,
  "best_seller": true,
  "prices": {
    "usd": 19.95,
    "gbp": 17.95,
    "eur": 18.95
  }
}
```



Indexing a document with ID 2

```
PUT books/_doc/3
{
  "title": "Java: A Beginner's Guide",
  "author": "Herbert Schildt",
  "release_date": "2018-11-20",
  "amazon_rating": 4.2,
  "best_seller": true,
  "prices": {
    "usd": 19.99,
    "gbp": 19.99,
    "eur": 19.99
  }
}
```



Indexing a document with ID 3

# The JSON response for a \_count API invocation


GET books/\_count



```
{  
  "count" : 3,  
  "_shards" : {  
    "total" : 1,  
    "successful" : 1,  
    "skipped" : 0,  
    "failed" : 0  
  }  
}
```

The count variable indicates the total number of documents

# Fetching a book document by an ID



# Retrieving documents given a set of IDs

A GET request on a `_search` endpoint with a query to fetch the multiple documents

GET `books/_search`

{

`"query": {`

`"ids": {`

`"values": [1,2,3]`

`}`

}

The `ids` query expects an array with document IDs


The response returns all three documents

```
"hits" : [{  
  "_index" : "books",  
  "_type" : "_doc",  
  "_id" : "1",  
  "_source" : {  
    "title" : "Effective Java",  
    ...  
  }}, {  
  "_index" : "books",  
  "_id" : "2",  
  ...  
}, ...]
```

# Retrieving all documents using the search API

The generic search (with no request body) will return all the books from the index

GET books/\_search



The actual book is enclosed in the `_source` object

```
{  
  ...  
  "hits" : {  
    "total" : {  
      "value" : 3,  
      "relation" : "eq"  
    },  
    "max_score" : 1.0,  
    "hits" : [  
      {  
        "_index" : "books",  
        "_type" : "_doc",  
        "_id" : "1",  
        "_score" : 1.0,  
        "_source" : {  
          "title" : "Effective Java",  
          "author" : "Joshua Bloch"  
        }  
      },  
      ...  
    ]  
  }  
}
```

The `hits` object indicates the number of results returned

The `hits` array is composed of the actual returned results

# Fetching books

A match query fetching all books written by Joshua


```
GET books/_search
{
  "query": {
    "match": {
      "author": "Joshua"
    }
  }
}
```

The match query with an author clause

The response returns document with a match.

```
"hits" : [
  {
    "_index" : "books",
    "_type" : "_doc",
    "_id" : "1",
    "_score" : 1.0417082,
    "_source" : {
      "title" : "Effective Java",
      "author" : "Joshua Bloch"
      ...
    }
  }
]
```

# Elasticsearch with data



# JSON vs. a relational database table structure

JSON representation of a Student data in Elasticsearch (the data is denormalized)

```
{  
    "title": "John Doe",  
    "date_of_birth": "1972-14-03",  
    "age": 23,  
    "address": {  
        //...  
    }  
}
```

Records are split into two individual tables and joined up by a foreign key between these two tables

Student data represented as relational data in a database (the data is normalized)

STUDENT TABLE				
ID	TITLE	DATE_OF_BIRTH	AGE	ADDRESS_ID
1	John Doe	1972-14-03	23	123456

ADDRESS TABLE		
ADDRESS_ID	ADDRESS_LINE1	POSTCODE
123456	34, Johndoe Land, London	LD1DNN

# Removing document types

The url consists of the type (car) of the document too (which is deprecated and be removed in version 8.0)

```
PUT cars/car/1
{
  "make": "Toyota",
  "model": "Avensis"
}
```

> V7.0: The explicit type is replaced with an endpoint named `_doc` (not document type)


```
PUT cars/_doc/1
{
  ...
}
```

```
#! [types removal] Specifying types in
document index requests is deprecated, use the
typeless endpoints instead
({index}/_doc/{id}, /{index}/_doc, or
/{index}/_create/{id}).
```


```
{
  "_index" : "cars",
  "_type" : "car",
  "_id" : "1",
  "_version" : 1,
  "result" : "created",
  "_shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "_seq_no" : 0,
  "_primary_term" : 1
}
```

While using the type upto 7.x is allowed (you will receive a warning as shown here), it is advisable to drop the type completely

# An index




# Lucene's mechanism for indexing documents



# Engine is not ready showing RED status

Node A has two primary shards ready  
and third shard is being instantiated

The shard 3 has not  
yet been instantiated





All shards are not yet assigned.  
Cluster status is RED

# A single node with three shards


Node A has three primary shards,  
all instantiated and ready

All shards are  
ready for action




All shards are ready but replicas are not  
yet assigned. Cluster status is YELLOW


# Shards balanced on new node but replicas not assigned




# All shards and replicas are allocated




# Health of shards using a traffic light signal board




# Replicas were lost when a node crashed



# A single node Elasticsearch cluster



# From a single node to a multiple node



# An inverted index data structure

**Inverted Index**

Two full-text fields:  
“Hello, World!”  
“Hello, Mate”



Word	Doc Num
hello	1, 2
world	1
mate	2

# Relevant results for Java in a title search

```
GET books/_search
{
  "_source": "title",
  "query": {
    "match": {
      "title": "Java"
    }
  }
}
```



```
"hits" : [
  ...
  "max_score" : 0.33537668,
  ...
  "hits" : [
    {
      "_score" : 0.33537668,
      "_source" : { "title" : "Effective Java" }
    },
    {
      "_score" : 0.30060259,
      "_source" : { "title" : "Head First Java" }
    },
    {
      "_score" : 0.18531466,
      "_source" : { "title" : "Test-Driven: TDD and Acceptance TDD for Java Developers" }
    }...
  ]
}
```